

Assignment_5

Niharika

2024-04-07

```
library(class)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
## Loading required package: lattice
```

```
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 4.3.3
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.3.3
```

```
## Warning: package 'readr' was built under R version 4.3.3
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
## Warning: package 'forcats' was built under R version 4.3.3
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats   1.0.0      ✓ stringr   1.5.1
## ✓ lubridate 1.9.3      ✓ tibble    3.2.1
## ✓ purrr     1.0.2      ✓ tidyr     1.3.1
```

```
## — Conflicts ————— tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## ✗ purrr::lift()    masks caret::lift()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(ISLR)
library(factoextra)
```

```
## Warning: package 'factoextra' was built under R version 4.3.3
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(dbscan)
```

```
## Warning: package 'dbscan' was built under R version 4.3.3
```

```
##
## Attaching package: 'dbscan'
##
## The following object is masked from 'package:stats':
##
##      as.dendrogram
```

```
library(cluster)
```

```
## Warning: package 'cluster' was built under R version 4.3.3
```

```
library(klustR)
```

```
## Warning: package 'klustR' was built under R version 4.3.3
```

```
library(ggplot2)
library(dplyr)
library(gridExtra)
```

```
## Warning: package 'gridExtra' was built under R version 4.3.3
```

```
##
## Attaching package: 'gridExtra'
##
## The following object is masked from 'package:dplyr':
##
##      combine
```

```
cereals.data <- read.csv("C:/Users/nihar/Downloads/Cereals.csv")

# Omit the null values
cereals.data <- na.omit(cereals.data)

# check the dimensions of the dataset
dim(cereals.data)
```

```
## [1] 74 16
```

```
head(cereals.data)
```

```
##           name mfr type calories protein fat sodium fiber carbo
## 1      100%_Bran  N   C       70        4   1   130  10.0   5.0
## 2    100%_Natural_Bran  Q   C      120        3   5    15   2.0   8.0
## 3         All-Bran  K   C       70        4   1   260   9.0   7.0
## 4 All-Bran_with_Extra_Fiber  K   C       50        4   0   140  14.0   8.0
## 6  Apple_Cinnamon_Cheerios  G   C      110        2   2   180   1.5  10.5
## 7        Apple_Jacks  K   C      110        2   0   125   1.0  11.0
##  sugars potass vitamins shelf weight cups  rating
## 1      6    280      25    3      1 0.33 68.40297
## 2      8    135       0    3      1 1.00 33.98368
## 3      5    320      25    3      1 0.33 59.42551
## 4      0    330      25    3      1 0.50 93.70491
## 6     10     70      25    1      1 0.75 29.50954
## 7     14     30      25    2      1 1.00 33.17409
```

```
# To know the column names of the dataset
t(t(names(cereals.data)))# The 't' function creates a transpose of the dataframe
```

```
##      [,1]
## [1,] "name"
## [2,] "mfr"
## [3,] "type"
## [4,] "calories"
## [5,] "protein"
## [6,] "fat"
## [7,] "sodium"
## [8,] "fiber"
## [9,] "carbo"
## [10,] "sugars"
## [11,] "potass"
## [12,] "vitamins"
## [13,] "shelf"
## [14,] "weight"
## [15,] "cups"
## [16,] "rating"
```

Assignment Task A #“Apply hierarchical clustering to the data using Euclidean distance to the normalized measurements. Use Agnes to compare the clustering from single linkage, complete linkage, average linkage, and Ward. Choose the best method.”

```
clust.data <- cereals.data[,4:16]
dim(clust.data)
```

```
## [1] 74 13
```

```
head(clust.data)
```

```
##   calories protein fat sodium fiber carbo sugars potass vitamins shelf weight
## 1      70      4  1   130  10.0   5.0     6   280      25     3     1
## 2     120      3  5    15   2.0   8.0     8   135       0     3     1
## 3      70      4  1   260   9.0   7.0     5   320      25     3     1
## 4      50      4  0   140  14.0   8.0     0   330      25     3     1
## 6     110      2  2   180   1.5  10.5    10    70      25     1     1
## 7     110      2  0   125   1.0  11.0    14    30      25     2     1
##   cups   rating
## 1 0.33 68.40297
## 2 1.00 33.98368
## 3 0.33 59.42551
## 4 0.50 93.70491
## 6 0.75 29.50954
## 7 1.00 33.17409
```

```
summary(clust.data)
```

```
##      calories      protein      fat      sodium      fiber
## Min.   : 50      Min.   :1.000      Min.   :0      Min.   : 0.0      Min.   : 0.000
## 1st Qu.:100      1st Qu.:2.000      1st Qu.:0      1st Qu.:135.0      1st Qu.: 0.250
## Median :110      Median :2.500      Median :1      Median :180.0      Median : 2.000
## Mean   :107      Mean   :2.514      Mean   :1      Mean   :162.4      Mean   : 2.176
## 3rd Qu.:110      3rd Qu.:3.000      3rd Qu.:1      3rd Qu.:217.5      3rd Qu.: 3.000
## Max.   :160      Max.   :6.000      Max.   :5      Max.   :320.0      Max.   :14.000
##      carbo      sugars      potass      vitamins
## Min.   : 5.00      Min.   : 0.000      Min.   : 15.00      Min.   : 0.00
## 1st Qu.:12.00      1st Qu.: 3.000      1st Qu.: 41.25      1st Qu.: 25.00
## Median :14.50      Median : 7.000      Median : 90.00      Median : 25.00
## Mean   :14.73      Mean   : 7.108      Mean   : 98.51      Mean   : 29.05
## 3rd Qu.:17.00      3rd Qu.:11.000      3rd Qu.:120.00      3rd Qu.: 25.00
## Max.   :23.00      Max.   :15.000      Max.   :330.00      Max.   :100.00
##      shelf      weight      cups      rating
## Min.   :1.000      Min.   :0.500      Min.   :0.2500      Min.   :18.04
## 1st Qu.:1.250      1st Qu.:1.000      1st Qu.:0.6700      1st Qu.:32.45
## Median :2.000      Median :1.000      Median :0.7500      Median :40.25
## Mean   :2.216      Mean   :1.031      Mean   :0.8216      Mean   :42.37
## 3rd Qu.:3.000      3rd Qu.:1.000      3rd Qu.:1.0000      3rd Qu.:50.52
## Max.   :3.000      Max.   :1.500      Max.   :1.5000      Max.   :93.70
```


```
scaled.data <- scale(clust.data)
head(scaled.data)
```

```
##      calories      protein      fat      sodium      fiber      carbo      sugars
## 1 -1.8659155  1.3817478  0.0000000 -0.3910227  3.22866747 -2.5001396 -0.2542051
## 2  0.6537514  0.4522084  3.9728810 -1.7804186 -0.07249167 -1.7292632  0.2046041
## 3 -1.8659155  1.3817478  0.0000000  1.1795987  2.81602258 -1.9862220 -0.4836096
## 4 -2.8737823  1.3817478 -0.9932203 -0.2702057  4.87924705 -1.7292632 -1.6306324
## 6  0.1498180 -0.4773310  0.9932203  0.2130625 -0.27881412 -1.0868662  0.6634132
## 7  0.1498180 -0.4773310 -0.9932203 -0.4514312 -0.48513656 -0.9583868  1.5810314
##      potass      vitamins      shelf      weight      cups      rating
## 1  2.5605229 -0.1818422  0.9419715 -0.2008324 -2.0856582  1.8549038
## 2  0.5147738 -1.3032024  0.9419715 -0.2008324  0.7567534 -0.5977113
## 3  3.1248675 -0.1818422  0.9419715 -0.2008324 -2.0856582  1.2151965
## 4  3.2659536 -0.1818422  0.9419715 -0.2008324 -1.3644493  3.6578436
## 6 -0.4022862 -0.1818422 -1.4616799 -0.2008324 -0.3038480 -0.9165248
## 7 -0.9666308 -0.1818422 -0.2598542 -0.2008324  0.7567534 -0.6553998
```

```
dim(scaled.data)
```

```
## [1] 74 13
```

```
distance <- get_dist(scaled.data)
fviz_dist(distance)
```



```
# Consider the Euclidean distance to know the distance between each point
dist <- dist(scaled.data, method = "euclidean")
```

```
# cluster using different Linkage methods
```

```
# Single linkage
```

```
hc_single <- agnes(dist,method = "single")
```

```
# Complete Linkage
```

```
hc_complete <- agnes(dist,method = "complete")
```

```
# average linkage
```

```
hc_average <- agnes(dist,method = "average")
```

```
# Ward's Linkage
```

```
hc_ward <- agnes(dist,method = "ward")
```

```
#comparing the above Links values with Agglomerative coefficient(ac)
```

```
print(hc_single$ac)
```

```
## [1] 0.6067859
```

```
print(hc_complete$ac)
```

```
## [1] 0.8353712
```

```
print(hc_average$ac)
```

```
## [1] 0.7766075
```

```
print(hc_ward$ac)
```

```
## [1] 0.9046042
```

From the output we can see that the agglomerative coefficient for wards method is high. so we need to consider that.


Assignment Task B #“How many clusters would you choose?” I choose (3-8)

```
# Cluster the data using Hierarchical clustering
hc3 <- hclust(dist, method = "ward.D")

# Plot the dendrogram
plot(hc3, cex = 0.1, main = "Dendrogram of Hierarchical Clustering for No.of clusters=3")

# Split the tree
data3 = rect.hclust(hc3, k=3 , border = 1:3)
```

Dendrogram of Hierarchical Clustering for No.of clusters=3



```
clusters3 <- cutree(hc3, k=3)


# To know the cluster size
data_with_clusters3 <- cbind(Cluster = data3)
data_with_clusters3
```

```
##      Cluster
## [1,] integer,21
## [2,] integer,18
## [3,] integer,35
```

```
# add the clusters to the table
clustered <- cbind(cereals.data, clusters3)


# heatmap for the data
heatmap(as.matrix(scaled.data), Colv = NA, hclustfun = hclust,
        col=rev(paste("gray",1:99,sep="")))

```

```
# plot the data
```

```
ggplot(clustered, aes(factor(clusters3), fill= mfr))+ geom_bar(position='dodge')+labs(x='clusters3')
```




```
# Cluster the data using Hierarchical clustering
hc4 <- hclust(dist, method = "ward.D")

# Plot the dendrogram
plot(hc4, cex = 0.1, main = "Dendrogram of Hierarchical Clustering for No.of clusters=4")

# Split the tree
data4 = rect.hclust(hc4, k=4 , border = 1:3)
```

Dendrogram of Hierarchical Clustering for No.of clusters=4




```
clusters4 <- cutree(hc4, k=4)

# To know the cluster size
data_with_clusters <- cbind(Cluster = data4)
data_with_clusters
```

```
##      Cluster
## [1,] integer,21
## [2,] integer,18
## [3,] integer,23
## [4,] integer,12
```

```
# add the clusters to the table
clustered <- cbind(cereals.data, clusters4)

# plot the graph
ggplot(clustered, aes(factor(clusters4), fill= mfr))+ geom_bar(position='dodge')+labs(x='clusters4')
```



Assignment C #“Comment on the structure of the clusters and on their stability. Hint: To check stability, partition the data and see how well clusters formed based on one part apply to the other part. To do this:

Ans.To check stability of clusters, the data set will be split into a 70/30 partition. The 70% will be used to create cluster assignments again, and then the remaining 30% will be assigned based on their closest centroid.

```
set.seed(111780)
# Split the data into 70% partition A and 30% partition B
cerealIndex <- createDataPartition(cereals.data$protein, p=0.3, list =
F)
cereal_preprocessed_PartitionB <- cereals.data[cerealIndex, ]
cereal_preprocessed_PartitionA <- cereals.data[-cerealIndex,]
```

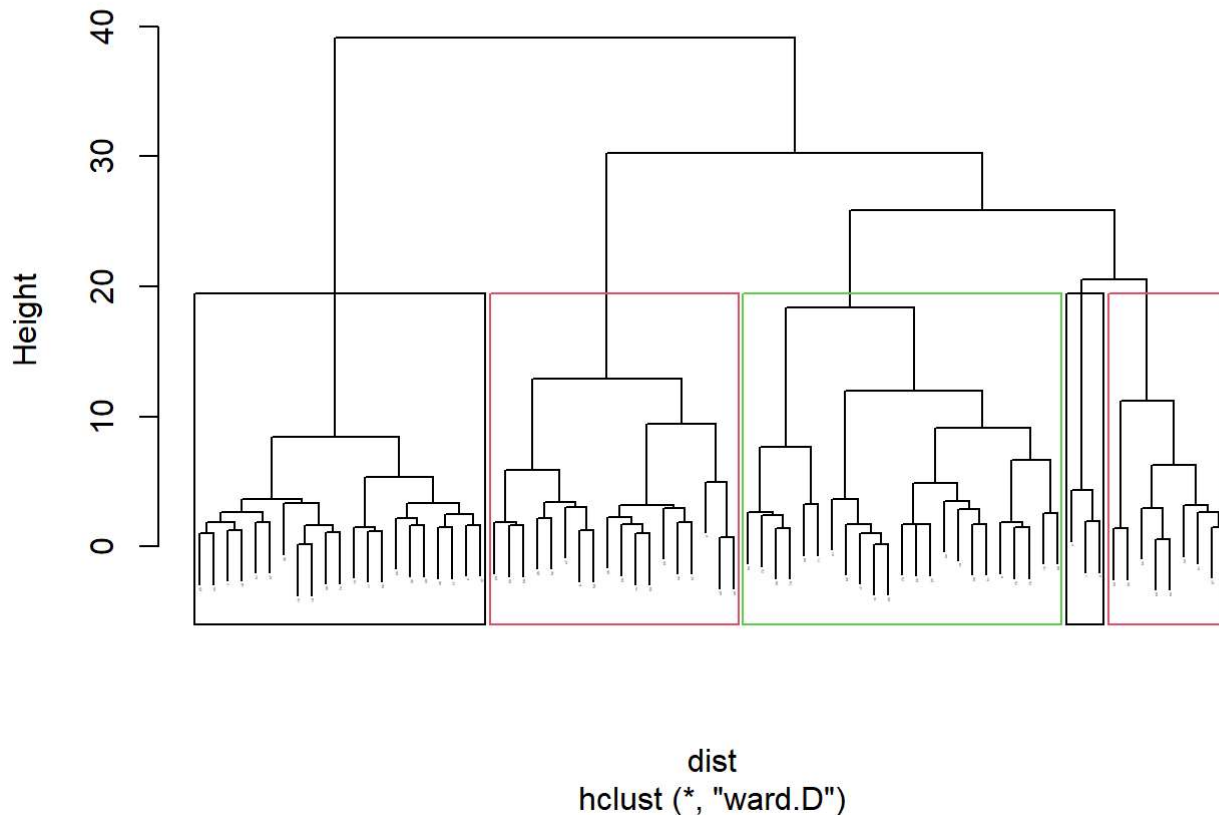
#1. Cluster partition A #2. Use the cluster centroids from A to assign each record in partition B (each record is assigned to the cluster with the closest centroid). #3. Assess how consistent the cluster assignments are compared to the assignments based on all the data”

```
# Cluster the data using Hierarchical clustering
hc5 <- hclust(dist, method = "ward.D")

# Plot the dendrogram
plot(hc5, cex = 0.1, main = "Dendrogram of Hierarchical Clustering for No.of clusters=5")

# Split the tree
data5 = rect.hclust(hc5, k=5 , border = 1:3)
```

Dendrogram of Hierarchical Clustering for No.of clusters=5



```
clusters5 <- cutree(hc5, k=5)


# To know the cluster size
data_with_clusters <- cbind(Cluster = data5)
data_with_clusters
```

```
##      Cluster
## [1,] integer,21
## [2,] integer,18
## [3,] integer,23
## [4,] integer,3
## [5,] integer,9
```

```
# add the data
clustered <- cbind(cereals.data, clusters5)


# heatmap for the data
heatmap(as.matrix(scaled.data), Colv = NA, hclustfun = hclust,
        col=rev(paste("gray",1:99,sep="")))

```



```
# plot the data
```

```
ggplot(clustered, aes(factor(clusters5), fill= mfr))+ geom_bar(position='dodge')+labs(x='clusters5')
```




```
# Cluster the data using Hierarchical clustering
hc6 <- hclust(dist, method = "ward.D")

# Plot the dendrogram
plot(hc6, cex = 0.1, main = "Dendrogram of Hierarchical Clustering for No.of clusters=6")

# Split the tree
data6 = rect.hclust(hc6, k=6 , border = 1:3)
```

Dendrogram of Hierarchical Clustering for No.of clusters=6



```
clusters6 <- cutree(hc6, k=6)


# To know the cluster size
data_with_clusters <- cbind(Cluster = data6)
data_with_clusters
```

```
##      Cluster
## [1,] integer,21
## [2,] integer,18
## [3,] integer,6
## [4,] integer,17
## [5,] integer,3
## [6,] integer,9
```

```
# add the data
clustered <- cbind(cereals.data, clusters6)


# heatmap for the data
heatmap(as.matrix(scaled.data), Colv = NA, hclustfun = hclust,
        col=rev(paste("gray",1:99,sep="")))

```

```
# plot the data
```

```
ggplot(clustered, aes(factor(clusters6), fill= mfr))+ geom_bar(position='dodge')+labs(x='clusters6')
```




```
# Cluster the data using Hierarchical clustering
hc8 <- hclust(dist, method = "ward.D")

# Plot the dendrogram
plot(hc8, cex = 0.1, main = "Dendrogram of Hierarchical Clustering for No.of clusters=8")

# Split the tree
data8 = rect.hclust(hc8, k=8 , border = 1:3)
```

Dendrogram of Hierarchical Clustering for No.of clusters=8



```
clusters8 <- cutree(hc8, k=8)


# To know the cluster size
data_with_clusters <- cbind(Cluster = data8)
data_with_clusters
```

```
##      Cluster
## [1,] integer,21
## [2,] integer,8
## [3,] integer,10
## [4,] integer,6
## [5,] integer,5
## [6,] integer,12
## [7,] integer,3
## [8,] integer,9
```

```
# add the data
clustered <- cbind(cereals.data, clusters8)


# heatmap for the data
heatmap(as.matrix(scaled.data), Colv = NA, hclustfun = hclust,
        col=rev(paste("gray",1:99,sep="")))

```



```
# plot the data
```

```
ggplot(clustered, aes(factor(clusters8), fill= mfr))+ geom_bar(position='dodge')+labs(x='clusters8')
```



Assignment Task D

In the context of selecting cereals for elementary school cafeterias to ensure a healthy diet, it's important to consider the normalization of data. Normalizing the data may not be appropriate because it would scale the nutritional information based on the sample of cereals being analyzed. This approach could be misleading because the dataset might include cereals with extremely high sugar content and very low fiber, iron, and other essential nutrients. When normalized across the sample set, it becomes difficult to interpret the nutritional value accurately. For instance, a cereal with a normalized value of 0.999 for iron might appear to have nearly all the nutritional iron a child needs, but it could simply be the best among a set of unhealthy options. A more suitable approach for preprocessing the data would be to express it as a ratio to the daily recommended intake of calories, fiber, carbohydrates, etc., for a child. This method allows analysts to make better-informed decisions about clustering without allowing a few variables with larger values to dominate the distance calculations. By reviewing the clusters, analysts can assess the average nutritional values within each cluster to determine what percentage of a student's daily recommended nutrition would come from each cereal. This approach enables staff to make informed decisions about selecting "healthy" cereals for the cafeteria.